

Goals:

- Apply ideas from Union-Find
- Explore more heap variants and their properties
- Reinforce algorithm design principles

sorted Array
cost of looking
Heap is $O(n)$ NOT $O(n \log n)$

Problem 1. A whole lot of work to do...

Suppose you have a huge amount of work to do. To keep track of the n tasks, you have a big array W where $W[i] = 0$ means that task i has not yet been completed, and $W[i] = 1$ implies that task i has been completed.

0	0	1	0	0	1	1	1	1	1	0	0
0	1	2	3	4	5	6	7	8	9	10	11

Figure 1: An example of W .

In order to help coordinate getting all the work done, you want to implement a data structure that implements the following three operations:

Operation	Behaviour
<code>lookup(i)</code>	Returns the value $W[i]$.
<code>mark(i)</code>	Marks task i as completed, i.e., sets $W[i] \leftarrow 1$.
<code>nextTask(i)</code>	Returns the next task from i onwards that is not yet completed, i.e., the next index $j \geq i$ where $W[j] = 0$.

For example, for the array above (Figure 1), `lookup(2)` returns 1, `mark(4)` would change the zero to a one, and `nextTask(6)` would return 10.

The simplest solution, of course, is just to use the array. In that case, it is easy to implement the `lookup` and `mark` operations in $O(1)$ time, but the `nextTask` operation may take $\Omega(n)$ time.

The goal in this problem is to develop data structures that provide different trade-offs in performance. For all versions, we want `lookup` to complete in $O(1)$ time.

Problem 1.a. Design a data structure where the `mark` and `nextTask` operations complete in $O(\log n)$ time, worst-case.

Problem 1.b. Design a data structure where the **mark** operation runs in $O(\log n)$ *amortized* time and **nextTask** operation runs in $O(1)$ time, worst-case.

Problem 1.c. Design a data structure where the **mark** operation runs in worst-case $O(1)$ time and **nextTask** operation runs in $O(\alpha(n))$ *amortized* time (where $\alpha(n)$ refers to the time complexity of the inverse Ackermann function, i.e., the amortized time for executing n operations on a union-find data structure).

Problem 2. Communist Data Structures

In this problem we will build a new type of mergeable Max-Heap. It is often useful to be able to merge data structures. For example, they can be used to build divide-and-conquer algorithms, they can be used more easily in augmenting trees, etc.

Let us first define the following terms:

Term	Definition
right spine	The sequence of nodes traversed in a tree if you start at a node and always go right until you find a node with no right child (which may not be a leaf).
$u.\text{rightRank}$	The number of nodes along the right spine of node u .
LEFTIST property	The property a tree satisfies if, for every node, $\text{rightRank}(L) \geq \text{rightRank}(R)$ where L and R are the left and right child respectively. The rightRank for a non-existent child is treated to be zero.
LEFTIST (Max) HEAP	A tree that satisfies both the (Max) Heap order property and the LEFTIST property.

Below is the ADT specification of a LEFTIST HEAP.

Operation	Behaviour
insert (u)	Insert node u into the heap.
merge ($t2$)	Returns the tree as a result of merging the heap with $t2$.
extractMax ()	Removes the node with the maximum key from the heap and return it.




Figure 2: Example of a LEFTIST HEAP, where the `rightRank` for each node is labelled next to it.

Problem 2.a. What is the maximum height/depth of a LEFTIST HEAP?

Problem 2.b. What is the maximum `rightRank` of a LEFTIST HEAP containing n nodes?

Problem 2.c. How would you implement the `insert` operation for a LEFTIST HEAP?

Where would be a good place to insert the new node? What is the maximum cost?

Note that your strategy must work on *any valid* LEFTIST HEAP. i.e. it need not be one built sequentially from the insertion algorithm you propose.

Hints: Think recursively! Where is an efficient place to insert it? Realise also you can always swap the left and right subtrees of a node if it does not satisfy the LEFTIST property.

Problem 2.d. Now come up with an algorithm to implement the `merge` operation. What is the running time?

Hint: Apply the same idea from part b.

Problem 2.e. How would you implement `extractMax` in a LEFTIST HEAP?

Problem 2.f. Now suppose we extend the LEFTIST HEAP ADT with the following additional operations:

Operation	Behaviour
<code>delete(u)</code>	Delete the node u from the heap.
<code>updateKey(u, k)</code>	Update the key of node u in the heap to k .

How would you implement them efficiently? Note that for both these operations, you do not have to search for node u since it is provided as an argument.

Hint: You should make use of what you have already come up with from earlier parts.